DealReasoner

Verifiable reasoning for freelancer client go / no-go decisions.

Every claim cites its evidence. When evidence is thin,

it escalates instead of guessing.

github.com/Shivay00001/dealreasoner

The problem

Freelancers get vague inbound briefs every week.

Each discovery call costs an hour — most were never a fit.

No structured way to turn a vague brief into a reliable decision.

DealReasoner turns a brief into GO / PILOT / NO-GO / ESCALATE.

How it reasons

ACT — gather evidence: brief signals, deal memory, profile constraints.

Each evidence item gets a stable ID (E1, E2, ...).

THINK — propose 3 competing hypotheses, each citing evidence IDs.

TRACE VERIFY — drop any claim that cites no valid evidence.

DECIDE — score deal-fit, calibrate the verdict, apply policy.

Pipeline: brief -> ACT -> THINK -> TRACE VERIFY -> DECIDE -> verdict

Live demo — GO case (real output)

Brief: Dubai real estate agency, 40 agents, $5k setup budget,

WhatsApp lead-qualification agent.

VERDICT: GO · confidence 0.89 · scorer: Jev

Why: Deal-fit 0.77 via Jev over 9 evidence items [E1].

Reasoning trace: 3 cited claims kept, 0 uncited dropped.

Ask on the call: success in 30 days? deadline? prior vendor?

Live demo — ESCALATE case (real output)

Brief: clinic chain exploring ad management; no budget approved yet.

VERDICT: ESCALATE · confidence 0.73 · scorer: Jev

Why: Not enough evidence to decide [E1]: budget undisclosed

and no red flags to rule on — human review needed.

This is the reliability behaviour: abstain instead of guessing.

Measured, not claimed

Eval set: 12 briefs with expected verdicts (src/eval/eval_set.json).

Latest runs: 12/12 correct. One earlier run: 11/12 —

live Jev scoring varies slightly near thresholds.

Abstention policy: thin evidence or low scorer confidence

-> ESCALATE to the human. Pilot-first cap: cold deals

at or under $2,000 start as PILOT, never GO.

Reproduce: python -m src.eval.run_eval (stdlib only, $0)

Sponsor tech — honest status

Zetaris (data layer): client implemented + tested against the documented REST shape; live Cloud dataset + token pending console access. Local-KB fallback is labeled in output.

Nemotron via Meterless: client implemented + tested (planner model for THINK, fast model for responder); API key issued in the HackOS workspace.

Cursor: Agent CLI installed (2026.10.22-27dcf5f); cloud runs need a Cursor API key (login step, $0 budget — not available here).

Jev scoring: LIVE in every run shown.

Why this wins Track 03

Reasoning architecture, not a chatbot wrapper: evidence IDs,

competing hypotheses, verified trace, calibrated abstention.

Impact: saves freelancers hours of dead discovery calls,

with a decision they can audit line by line.

Technical: stdlib-only, reproducible in one command, $0 to run.

Demo: real terminal output, GO + ESCALATE cases, 126-second video.

github.com/Shivay00001/dealreasoner